

# Modeling and Analysis of Time Series Data

## Chapter 11: Introduction to partially observed Markov process models

STATS 531, Winter 2026

Edward Ionides

## Stochastic dynamic systems observed with noise

- ▶ Uncertainty and variability are common features of biological and social systems. Complex physical systems can also be unpredictable: we can only forecast weather reliably in the near future.
- ▶ Time series models of deterministic trend plus colored noise imply perfect predictability if the trend function enables extrapolation.
- ▶ To model variability and unpredictability in a dynamic system, we can specify a stochastic (i.e., random) model for the system.
- ▶ Often times, the full dynamic system is unobserved. We have only noisy or incomplete measurements.
- ▶ We model measurements as random variables conditional on the trajectory of the *latent process*. The latent process is also called a *state process* or *hidden process*.

## The Markov property

- ▶ A model for a stochastic dynamic system has the *Markov property* if the future evolution of the system depends only on the current state, plus randomness introduced in future.
- ▶ A model with the Markov property may be called a *Markov chain* or a *Markov process*.
- ▶ We use the term Markov process since the term chain is often reserved for situations where either time or the latent state (or both) take discrete values.
- ▶ The Markov property is often used to model the latent process in a time series model.

## Notation for discrete time Markov processes

- ▶ A time series model  $X_{0:N}$  is a *Markov process* model if the conditional densities satisfy the *Markov property* [P1] that

$$[P1] \quad f_{X_n|X_{0:n-1}}(x_n | x_{0:n-1}) = f_{X_n|X_{n-1}}(x_n | x_{n-1})$$

for all  $n \in 1 : N$ .

- ▶ We may suppose there is an underlying continuous time,  $t$ , such that  $X_n$  occurs at time  $t_n$ .
- ▶ We write  $X(t)$  for the continuous time model, setting  $X_n = X(t_n)$ .
- ▶  $t_{1:N}$  are *measurement times*.
- ▶  $t_0$  is the *initialization time*.

## Initial conditions

- ▶ We *initialize* the Markov process model at a time  $t_0$ , although data are collected only at times  $t_{1:N}$ .
- ▶ The initialization model could be deterministic (a fixed value) or a random variable.
- ▶ We model  $X_0 = X(t_0)$  as a draw from a probability density function

$$f_{X_0}(x_0). \quad (1)$$

- ▶ A fixed initial value is a special case of a density corresponding to a point mass with probability one at the fixed value.
- ▶ A discrete probability mass function is a special case of a density corresponding to a collection of point masses.

## The process model

- ▶ The probability density function  $f_{X_n|X_{n-1}}(x_n | x_{n-1})$  is called the *one-step transition density* of the Markov process.
- ▶ The Markov property asserts that the next step taken by a Markov process follows the one-step transition density based on the current state, whatever the previous history of the process.
- ▶ For a Markov model, the full joint distribution of the latent process is entirely specified by the one-step transition densities, given the initial value.
- ▶ Therefore, we also call  $f_{X_n|X_{n-1}}(x_n | x_{n-1})$  the *process model*.

## The joint distribution in terms of one-step transition densities

**Exercise.** Use [P1] to derive an expression for the joint distribution of a Markov process as a product of the one-step transition densities. In other words, derive

$$[P2] \quad f_{X_{0:N}}(x_{0:N}) = f_{X_0}(x_0) \prod_{n=1}^N f_{X_n|X_{n-1}}(x_n | x_{n-1}).$$

**Hint:** This involves elementary rules for manipulation of joint and conditional densities, together with application of the Markov property. It is a good exercise to work through by hand to build familiarity with the model class.

**Question.** Explain why a causal Gaussian AR(1) process is a Markov process.

## Time-homogeneous transitions and stationarity

- ▶ The one-step transition density  $f_{X_n|X_{n-1}}$  for a Markov process  $X_{0:N}$  can depend on  $n$ .
- ▶  $X_{0:N}$  is *time-homogeneous* if  $f_{X_n|X_{n-1}}$  does not depend on  $n$ , so there is a conditional density  $f(\cdot | \cdot)$  such that, for all  $n \in 1:N$ ,

$$f_{X_n|X_{n-1}}(x_n | x_{n-1}) = f(x_n | x_{n-1}). \quad (2)$$

**Question.** If  $X_{0:N}$  is strictly stationary, it is time-homogeneous.  
Why?

**Question.** Time-homogeneity does not necessarily imply stationarity. Find a counter-example.

## Partially observed Markov process (POMP) models

- ▶ *Partial observation* may mean either or both of (i) measurement noise; (ii) entirely unmeasured latent variables.
- ▶ These features are present in many systems.
- ▶ A *partially observed Markov process* (POMP) model is defined by putting together a Markov latent process model and a *measurement model*.
- ▶ POMP models are a general class, covering many models designed for specific applications.
- ▶ Statistical methods for this general class give us flexibility to develop specific POMP models appropriate to a range of applications.

## The measurement model

- ▶ The *measurement process* is a collection of random variables  $Y_{1:N}$  which models the data  $y_{1:N}$ .
- ▶  $Y_n$  is assumed to depend on the latent process only through its value  $X_n$  at the time of the measurement. Formally, this assumption is:

$$\begin{aligned}[P3] f_{Y_n|X_{0:N}, Y_{1:n-1}, Y_{n+1:N}}(y_n | x_{0:N}, y_{1:n-1}, y_{n+1:N}) \\ = f_{Y_n|X_n}(y_n | x_n) \end{aligned}$$

- ▶ We call  $f_{Y_n|X_n}(y_n | x_n)$  the *measurement model*.

## Time-homogeneous measurement models

- ▶ In general, the measurement model can depend on  $n$  or on any covariate time series.
- ▶ The measurement model is *time-homogeneous* if there is a conditional probability density function  $g(\cdot | \cdot)$  such that, for all  $n \in 1 : N$ ,

$$f_{Y_n|X_n}(y_n | x_n) = g(y_n | x_n). \quad (3)$$

- ▶ Time-inhomogeneous process and measurement models are sufficiently common that we benefit from the extra generality of writing  $f_{X_n|X_{n-1}}(x_n|x_{n-1})$  and  $f_{Y_n|X_n}(y_n|x_n)$  versus  $f(x_n|x_{n-1})$  and  $g(y_n|x_n)$ .

# Prediction, filtering, smoothing and likelihood

## Four basic calculations for working with POMP models

Many time series models in science, engineering and industry can be written as POMP models. A reason that POMP models form a useful tool for statistical work is that there are convenient recursive formulas to carry out four basic calculations:

1. Prediction
2. Filtering
3. Smoothing
4. Likelihood calculation

## Prediction

- ▶ One-step prediction (also called forecasting) of the latent process at time  $t_{n+1}$  given data up to time  $t_n$  involves finding

$$f_{X_{n+1}|Y_{1:n}}(x_{n+1} | y_{1:n}). \quad (4)$$

- ▶ We may want to predict more than one time step ahead. However, one-step prediction turns out to be closely related to computing the likelihood function, and therefore central to statistical inference.
- ▶ Our prediction is a conditional probability density, not a point estimate. In the context of forecasting, this is called a *probabilistic forecast*. What are the advantages of a probabilistic forecast over a point forecast? Are there any disadvantages?

## Filtering

- ▶ The *filtering* calculation at time  $t_n$  is to find the conditional distribution of the latent process  $X_n$  given data  $y_{1:n}$  available at time  $t_n$ .
- ▶ Filtering involves calculating

$$f_{X_n|Y_{1:n}}(x_n | y_{1:n}). \quad (5)$$

- ▶ This can be evaluated numerically or algebraically. We will see that Monte Carlo methods can be a good tool.
- ▶ The name “filtering” comes from the history of signal processing. A noisy received signal was passed through capacitors and resistors to construct a band pass filter estimating the source signal, just like an optical filter removes unwanted frequencies of light.

## Smoothing

- ▶ In the context of a POMP model, smoothing involves finding the conditional distribution of  $X_n$  given all the data,  $y_{1:N}$ .
- ▶ So, the smoothing calculation is to find

$$f_{X_n|Y_{1:N}}(x_n | y_{1:N}). \quad (6)$$

## The likelihood

- ▶ The likelihood is the joint density of  $Y_{1:N}$  evaluated at the data,

$$f_{Y_{1:N}}(y_{1:N}). \quad (7)$$

- ▶ The model may depend on a parameter vector  $\theta$ . We can include  $\theta$  in all the joint and conditional densities above. Then, the *likelihood function* is the likelihood viewed as a function of  $\theta$ . We write

$$\mathcal{L}(\theta) = f_{Y_{1:N}}(y_{1:N}; \theta) \quad (8)$$

- ▶ If we can compute  $\mathcal{L}(\theta)$  then we can perform numerical optimization to get a maximum likelihood estimate.
- ▶ Likelihood evaluation and maximization lets us compute profile likelihood confidence intervals, carry out likelihood ratio tests, and make AIC model comparisons.

## The prediction formula

- ▶ One-step prediction of the latent process at time  $t_n$  given data up to time  $t_{n-1}$  can be computed recursively in terms of the filtering problem at time  $t_{n-1}$ , via the *prediction formula* for  $n \in 1 : N$ ,

$$[\text{P4}] \quad f_{X_n|Y_{1:n-1}}(x_n | y_{1:n-1}) = \int f_{X_{n-1}|Y_{1:n-1}}(x_{n-1} | y_{1:n-1}) f_{X_n|X_{n-1}}(x_n | x_{n-1}) dx_{n-1}.$$

- ▶ For the case  $n = 1$ , we let  $1 : k$  be the empty set when  $k = 0$ , so that  $f_{X_0|Y_{1:0}}(x_0 | y_{1:0})$  means  $f_{X_0}(x_0)$ . In other words, the filter distribution at time  $t_0$  is the initial density for the latent process, since at time  $t_0$  we have no data to condition on.

**Exercise.** Derive [P4] using the definition of a POMP model with elementary properties of joint and conditional densities.

## Hints for deriving the recursion formulas

Any general identity holding for densities must also hold when we condition everything on a new variable.

**Example 1.** From

$$f_{XY}(x, y) = f_X(x) f_{Y|X}(y | x) \quad (9)$$

we can condition on  $Z$  to obtain

$$f_{XY|Z}(x, y | z) = f_{X|Z}(x | z) f_{Y|XZ}(y | x, z). \quad (10)$$

**Example 2.** The prediction formula is a special case of the identity

$$f_{X|Y}(x | y) = \int f_{XZ|Y}(x, z | y) dz. \quad (11)$$

**Example 3.** A conditional form of Bayes' identity is

$$f_{X|YZ}(x | y, z) = \frac{f_{Y|XZ}(y | x, z) f_{X|Z}(x | z)}{f_{Y|Z}(y | z)}. \quad (12)$$

## The filtering formula

Filtering at time  $t_n$  can be computed by combining the new information in the datapoint  $y_n$  with the calculation of the one-step prediction of the latent process at time  $t_n$  given data up to time  $t_{n-1}$ .

- ▶ This is carried out via the *filtering formula* for  $n \in 1 : N$ ,

$$[P5] \quad f_{X_n|Y_{1:n}}(x_n | y_{1:n}) = \frac{f_{X_n|Y_{1:n-1}}(x_n | y_{1:n-1}) f_{Y_n|X_n}(y_n | x_n)}{f_{Y_n|Y_{1:n-1}}(y_n | y_{1:n-1})}.$$

**Exercise.** Derive [P5] using the definition of a POMP model with elementary properties of joint and conditional densities.

The prediction and filtering formulas are *recursive*. If they can be computed for time  $t_n$  then they enable the computation at time  $t_{n+1}$ .

## The conditional likelihood formula

- ▶ The denominator in the filtering formula [P5] is the *conditional likelihood* of  $y_n$  given  $y_{1:n-1}$ .
- ▶ It can be computed in terms of the one-step prediction density, via the *conditional likelihood formula*,

$$[P6] \quad f_{Y_n|Y_{1:n-1}}(y_n | y_{1:n-1}) = \int f_{X_n|Y_{1:n-1}}(x_n | y_{1:n-1}) f_{Y_n|X_n}(y_n | x_n) dx_n.$$

- ▶ To make this formula work for  $n = 1$ , we take advantage of the convention that  $1 : k$  is the empty set when  $k = 0$ .

## Computation of the likelihood and log likelihood

- ▶ The likelihood of the entire dataset,  $y_{1:N}$  can be found from [P6], using the identity

$$f_{Y_{1:N}}(y_{1:N}) = \prod_{n=1}^N f_{Y_n|Y_{1:n-1}}(y_n | y_{1:n-1}). \quad (13)$$

- ▶ This equation uses the convention that  $1 : k$  is the empty set when  $k = 0$ , so the first term in the product is

$$f_{Y_1|Y_{1:0}}(y_1 | y_{1:0}) = f_{Y_1}(y_1) \quad (14)$$

- ▶ If our model has an unknown parameter  $\theta$  then this gives the *log likelihood function* as a sum of conditional log likelihoods,

$$\ell(\theta) = \log \mathcal{L}(\theta) = \log f_{Y_{1:N}}(y_{1:N}; \theta) = \sum_{n=1}^N \log f_{Y_n|Y_{1:n-1}}(y_n | y_{1:n-1}; \theta).$$

## The smoothing recursions

- ▶ Smoothing is less fundamental for likelihood-based inference than filtering and one-step prediction.
- ▶ Nevertheless, sometimes we want to compute the smoothing density, so we develop some necessary formulas.
- ▶ The filtering and prediction formulas are recursions forward in time: a solution at time  $t_{n-1}$  is used for the computation at  $t_n$ .
- ▶ For smoothing, we have *backwards smoothing recursion formulas*,

$$[\text{P7}] \quad f_{Y_{n:N}|X_n}(y_{n:N} | x_n) = f_{Y_n|X_n}(y_n | x_n) f_{Y_{n+1:N}|X_n}(y_{n+1:N} | x_n).$$

$$\begin{aligned} [\text{P8}] \quad f_{Y_{n+1:N}|X_n}(y_{n+1:N} | x_n) \\ = \int f_{Y_{n+1:N}|X_{n+1}}(y_{n+1:N} | x_{n+1}) f_{X_{n+1}|X_n}(x_{n+1} | x_n) dx_{n+1}. \end{aligned}$$

## Combining recursions to find the smoothing distribution

The forwards and backwards recursion formulas together allow us to compute the *smoothing formula*,

$$[\text{P9}] \quad f_{X_n|Y_{1:N}}(x_n | y_{1:N}) = \frac{f_{X_n|Y_{1:n-1}}(x_n | y_{1:n-1}) f_{Y_{n:N}|X_n}(y_{n:N} | x_n)}{f_{Y_{n:N}|Y_{1:n-1}}(y_{n:N} | y_{1:n-1})}.$$

**Exercise.** Show how [P7], [P8] and [P9] follow from the basic properties of conditional densities combined with the Markov property.

**Hint:** you can write the left hand side of [P9] as  $f_{X|YZ}$  with  $X = X_n$ ,  $Y = Y_{1:n-1}$ ,  $Z = Y_{n:N}$ .

## Linear Gaussian POMP (LG-POMP) models

Linear Gaussian partially observed Markov process (LG-POMP) models have many applications across science and engineering.

- ▶ Gaussian ARMA models are LG-POMP models. The POMP recursion formulas give a computationally efficient way to obtain the likelihood of a Gaussian ARMA model.
- ▶ Smoothing splines (including the Hodrick-Prescott filter, which is a smoothing spline) can be written as an LG-POMP model.
- ▶ The *Basic Structural Model* is an LG-POMP used for econometric forecasting. It models a stochastic trend, seasonality, and measurement error, in a framework with econometrically interpretable parameters. This is more interpretable than fitting SARIMA.

If an LG-POMP model is appropriate, you avoid Monte Carlo computations needed for inference in general nonlinear POMPs.

## The general LG-POMP model

Suppose the latent process,  $X_{0:N}$ , and the observation process  $\{Y_n\}$ , take vector values with dimension  $d_X$  and  $d_Y$ . A general mean zero LG-POMP model is specified by

- ▶ A sequence of  $d_X \times d_X$  matrices,  $\mathbf{A}_{1:N}$ ,
- ▶ A sequence of  $d_X \times d_X$  covariances,  $\{\Sigma_{X,n}, n \in 0 : N\}$ ,
- ▶ A sequence of  $d_Y \times d_X$  matrices,  $\mathbf{B}_{1:N}$ ,
- ▶ A sequence of  $d_Y \times d_Y$  covariances,  $\{\Sigma_{Y,n}, n \in 1 : N\}$ .

We initialize with  $X_0 \sim N[0, \Sigma_{X,0}]$  and then define the entire LG-POMP model by a recursion for  $n \in 1 : N$ ,

$$[\text{LG1}] \quad X_n = \mathbf{A}_n X_{n-1} + \epsilon_n, \quad \epsilon_n \sim N[0, \Sigma_{X,n}],$$

$$[\text{LG2}] \quad Y_n = \mathbf{B}_n X_n + \eta_n, \quad \eta_n \sim N[0, \Sigma_{Y,n}].$$

Often, but not always, we will have a *time-homogeneous* LG-POMP model, with  $\mathbf{A}_n = \mathbf{A}$ ,  $\mathbf{B}_n = \mathbf{B}$ ,  $\Sigma_{X,n} = \Sigma_X$ , and  $\Sigma_{Y,n} = \Sigma_Y$  for  $n \in 1 : N$ .

## The LG-POMP representation of a Gaussian ARMA

- ▶ Let  $\{Y_n\}$  be a Gaussian ARMA( $p, q$ ) model with noise process  $\omega_n \sim N[0, \sigma^2]$ , defined by

$$Y_n = \sum_{j=1}^p \phi_j Y_{n-j} + \omega_n + \sum_{k=1}^q \theta_k \omega_{n-k}. \quad (15)$$

- ▶ We look for a time-homogeneous LG-POMP defined by [LG1] and [LG2] where  $Y_n$  is the first component of  $X_n$  with no measurement error.
- ▶ To do this, we define  $d_X = r = \max(p, q + 1)$  and

$$\mathbf{B} = (1, 0, 0, \dots, 0), \quad (16)$$

$$\Sigma_Y = 0. \quad (17)$$

- ▶ We require  $\mathbf{A}$  and  $\Sigma_X$  such that  $Y_n$  satisfies the ARMA equation.

We state a solution and see if it works out. Consider

$$X_n = \begin{pmatrix} Y_n \\ \phi_2 Y_{n-1} + \cdots + \phi_r Y_{n-r+1} + \theta_1 \omega_n + \cdots + \theta_{r-1} \omega_{n-r+2} \\ \phi_3 Y_{n-1} + \cdots + \phi_r Y_{n-r+1} + \theta_2 \omega_n + \cdots + \theta_{r-1} \omega_{n-r+3} \\ \vdots \\ \phi_r Y_{n-1} + \theta_{r-1} \omega_n \end{pmatrix}$$

We check that the ARMA equation matches the matrix equation

$$X_n = \mathbf{A} X_{n-1} + \begin{pmatrix} 1 \\ \theta_1 \\ \theta_2 \\ \vdots \\ \theta_{r-1} \end{pmatrix} \omega_n, \text{ for } \mathbf{A} = \begin{pmatrix} \phi_1 & 1 & 0 & \cdots & 0 \\ \phi_2 & 0 & 1 & \ddots & \vdots \\ \vdots & \vdots & \ddots & \ddots & 0 \\ \phi_{r-1} & 0 & \cdots & 0 & 1 \\ \phi_r & 0 & \cdots & 0 & 0 \end{pmatrix}$$

This is a time-homogeneous LG-POMP, with  $\mathbf{A}$ ,  $\mathbf{B}$  and  $\Sigma_Y$  as above and

$$\Sigma_X = \sigma^2 (1, \theta_1, \theta_2, \dots, \theta_{r-1})^T (1, \theta_1, \theta_2, \dots, \theta_{r-1}). \quad (18)$$

## Different POMPs can give the same model for $Y_{1:N}$

- ▶ There are other LG-POMP representations giving rise to the same ARMA model.
- ▶ When only one component of a latent process is observed, any model giving rise to the same observed component is indistinguishable from the data.
- ▶ Here, the LG-POMP model has order  $d_X^2 = r^2 = \max(p, q + 1)^2$  parameters. The ARMA model has order  $r$  parameters, so we expect many ways to parameterize the ARMA model as a special case of the much larger LG-POMP model.
- ▶ This unidentifiability can also arise for non-Gaussian POMPs, but it is easier to see in the Gaussian case.

## The basic structural model

- ▶ The *basic structural model* was developed for econometric analysis.
- ▶ It decomposes an observable process  $Y_{1:N}$  as the sum of a *level* ( $L_n$ ), a *trend* ( $T_n$ ) describing the rate of change of the level, and a monthly *seasonal component* ( $S_n$ ).
- ▶ The model supposes that the level, trend and seasonality are perturbed with Gaussian white noise at each time point,

$$[BSM1] \quad Y_n = L_n + S_n + \epsilon_n$$

$$[BSM2] \quad L_n = L_{n-1} + T_{n-1} + \xi_n$$

$$[BSM3] \quad T_n = T_{n-1} + \zeta_n$$

$$[BSM4] \quad S_n = -\sum_{k=1}^{11} S_{n-k} + \eta_n$$

where  $\epsilon_n \sim N[0, \sigma_\epsilon^2]$ ,  $\xi_n \sim N[0, \sigma_\xi^2]$ ,  $\zeta_n \sim N[0, \sigma_\zeta^2]$ , and  $\eta_n \sim N[0, \sigma_\eta^2]$ .

## Two common special cases of the basic structural model

- ▶ The *local linear trend* model is the basic structural model without the seasonal component,  $\{S_n\}$ .
- ▶ The *local level model* is the basic structural model without either the seasonal component,  $\{S_n\}$ , or the trend component,  $\{T_n\}$ . The local level model is therefore a random walk observed with measurement error.

## Initial values for the basic structural model

- ▶ To complete the model, we need to specify initial values.
- ▶ We have an example of the common problem of failing to specify initial values: these are not explained in the documentation of the R implementation of the basic structural model, StructTS. We could go through the source code to find out what it does.
- ▶ Incidentally, ?StructTS does give some advice which resonates with our experience earlier in the course that optimization for ARMA models is often imperfect.

“Optimization of structural models is a lot harder than many of the references admit. For example, the ‘AirPassengers’ data are considered in Brockwell & Davis (1996): their solution appears to be a local maximum, but nowhere near as good a fit as that produced by ‘StructTS’. It is quite common to find fits with one or more variances zero, and this can include  $\sigma_{\epsilon}^2$ .”

## The basic structural model is an LG-POMP model

[BSM1–4] can be put in matrix form. Set

$$X_n = (L_n, T_n, S_n, S_{n-1}, S_{n-2}, \dots, S_{n-10})^T, \quad (19)$$

$$Y_n = (1, 0, 1, 0, 0, \dots, 0)X_n + \epsilon_n. \quad (20)$$

We can identify matrices  $\mathbf{A}$ ,  $\mathbf{B}$ ,  $\Sigma_X$  and  $\Sigma_Y$  giving a time-homogeneous LG-POMP representation [LG1, LG2] for the basic structural model.

The transition matrix  $\mathbf{A}$  combines the dynamics of level, trend, and seasonal components appropriately.

## Spline smoothing and its LG-POMP representation

- ▶ Spline smoothing is a standard method to smooth scatter plots and time plots. For example, `smooth.spline` and `hpfilter` in R.
- ▶ A *smoothing spline* for an equally spaced time series  $y_{1:N}$  collected at times  $t_{1:N}$  is the sequence  $x_{1:N}$  minimizing the *penalized sum of squares (PSS)*, defined as

$$[\text{SS1}] \quad \text{PSS}(x_{1:N}; \lambda) = \sum_{n=1}^N (y_n - x_n)^2 + \lambda \sum_{n=3}^N (\Delta^2 x_n)^2.$$

- ▶ The spline is defined for all times, but here we are only concerned with its value at the times  $t_{1:N}$ .
- ▶ Here,  $\Delta x_n = (1 - B)x_n = x_n - x_{n-1}$ .

- ▶ The *smoothing parameter*,  $\lambda$ , penalizes  $x_{1:N}$  to prevent the spline from interpolating the data.
- ▶ If  $\lambda = 0$ , the spline will go through each data point, i.e.,  $x_{1:N}$  will interpolate  $y_{1:N}$ .
- ▶ If  $\lambda = \infty$ , the spline will be the ordinary least squares regression fit,

$$x_n = \alpha + \beta n, \quad (21)$$

since  $\Delta^2(\alpha + \beta n) = 0$ .

- ▶ Now consider the linear Gaussian model,

$$\begin{aligned} [\text{SS2}] \quad X_n &= 2X_{n-1} - X_{n-2} + \epsilon_n, \quad \epsilon_n \sim \text{iid } N[0, \sigma^2/\lambda] \\ [\text{SS3}] \quad Y_n &= X_n + \eta_n, \quad \eta_n \sim \text{iid } N[0, \sigma^2] \end{aligned}$$

- ▶ Note that  $\Delta^2 X_n = \epsilon_n$ .
- ▶ We will show that [SS1] is equivalent to [SS2,SS3].

## Constructing a linear Gaussian POMP (LG-POMP) model matching [SS2] and [SS3]

**Question.**  $\{X_n, Y_n\}$  defined in [SS2] and [SS3] is not quite an LG-POMP model. However, we can use  $\{X_n\}$  and  $\{Y_n\}$  to build an LG-POMP model. How?

## Deriving the penalized spline from the LG-POMP

The joint density of  $X_{1:N}$  and  $Y_{1:N}$  in [SS2,SS3] is

$$f_{X_{1:N} Y_{1:N}}(x_{1:N}, y_{1:N}) = f_{X_{1:N}}(x_{1:N}) f_{Y_{1:N}|X_{1:N}}(y_{1:N} | x_{1:N}).$$

Taking logs we get

$$\log f_{X_{1:N} Y_{1:N}}(x_{1:N}, y_{1:N}) = \log f_{X_{1:N}}(x_{1:N}) + \log f_{Y_{1:N}|X_{1:N}}(y_{1:N} | x_{1:N}).$$

[SS2,SS3] tell us that  $\{\Delta^2 X_n, n \in 1 : N\}$  and  $\{Y_n - X_n, n \in 1 : N\}$  are independent  $N[0, \sigma^2/\lambda]$  and  $N[0, \sigma^2]$ .  
Thus,

$$\log f_{X_{1:N} Y_{1:N}}(x_{1:N}, y_{1:N}; \sigma, \lambda) = \frac{-1}{2\sigma^2} \left[ \sum_{n=1}^N (y_n - x_n)^2 - \lambda \sum_{n=3}^N (\Delta^2 x_n)^2 \right] + C.$$

Here,  $C$  depends on  $\sigma$  and  $\lambda$  but not on  $y_{1:N}$ .  $C$  depends on the initial terms  $x_0$  and  $x_{-1}$ , but we suppose these can be ignored, for example by modeling them with an improper uniform density.

- ▶ Comparing this with [SS1], we see that maximizing the density  $f_{X_{1:N}Y_{1:N}}(x_{1:N}, y_{1:N}; \sigma, \lambda)$  as a function of  $x_{1:N}$  is the same problem as finding the smoothing spline by minimizing the penalized sum of squares.
- ▶ For a Gaussian density, the mode (i.e., the maximum of the density) is equal to the expected value. Therefore, we have

$$\begin{aligned}
\operatorname{argmin}_{x_{1:N}} \text{PSS}(x_{1:N}; \lambda) &= \operatorname{argmax}_{x_{1:N}} f_{X_{1:N}Y_{1:N}}(x_{1:N}, y_{1:N}; \sigma, \lambda) \\
&= \operatorname{argmax}_{x_{1:N}} \frac{f_{X_{1:N}Y_{1:N}}(x_{1:N}, y_{1:N}; \sigma, \lambda)}{f_{Y_{1:N}}(y_{1:N}; \sigma, \lambda)} \\
&= \operatorname{argmax}_{x_{1:N}} f_{X_{1:N}|Y_{1:N}}(x_{1:N} \mid y_{1:N}; \sigma, \lambda) \\
&= \mathbb{E}[X_{1:N} \mid Y_{1:N} = y_{1:N}; \sigma, \lambda].
\end{aligned}$$

- ▶ Because a (conditional) normal distribution is characterized by its (conditional) mean and variance, the smoothing calculation for an LG-POMP model involves finding the conditional mean and variance of  $X_n$  given  $Y_{1:N} = y_{1:N}$ .
- ▶ We conclude that the smoothing problem for this LG-POMP model is the same as the spline smoothing problem defined by [SS1].
- ▶ If you have experience using smoothing splines, this connection may help you transfer that experience to POMP models.
- ▶ Once you have experience with POMP models, this connection helps you understand spline smoothers that are commonly used in many applications.
- ▶ For example, the smoothing parameter  $\lambda$  could be selected by maximum likelihood for the POMP model.

Why do we penalize by  $\sum_n (\Delta^2 X_n)^2$  when smoothing?

**Question.** We found that the smoothing spline corresponds to a particular choice of LG-POMP model given by [SS2, SS3]. Why do we choose that penalty, rather than the equivalent penalty from some other LG-POMP model?

**Note:** This LG-POMP model is sometimes reasonable, but presumably there are other occasions when a different LG-POMP model would lead to superior performance.

## The Kalman filter

- ▶ The *Kalman filter* is the name given to the prediction, filtering and smoothing formulas [P4–P9] for the linear Gaussian partially observed Markov process (LG-POMP) model.
- ▶ Linear Gaussian models have Gaussian conditional distributions.
- ▶ The integrals in the general POMP formulas can be found exactly for the Gaussian distribution, leading to linear algebra calculations of conditional means and variances.
- ▶ The R function `arima()` uses a Kalman filter to evaluate the likelihood of an ARMA model (or ARIMA, SARMA, SARIMA).

## Review of the multivariate normal distribution

- ▶ A random variable  $X$  taking values in  $\mathbb{R}^{d_X}$  is *multivariate normal* with mean  $\mu_X$  and variance  $\Sigma_X$  if we can write

$$X = \mathbf{H}Z + \mu_X, \quad (22)$$

where  $Z$  is a vector of  $d_X$  independent identically distributed  $N[0, 1]$  random variables and  $\mathbf{H}$  is a  $d_X \times d_X$  matrix square root of  $\Sigma_X$ , i.e.,

$$\mathbf{H}\mathbf{H}^T = \Sigma_X. \quad (23)$$

- ▶ A matrix square root of this type exists for any covariance matrix, though the choice of  $\mathbf{H}$  is not unique.
- ▶ We write  $X \sim N[\mu_X, \Sigma_X]$ . If  $\Sigma_X$  is invertible,  $X$  has a probability density function,

$$f_X(x) = \frac{1}{(2\pi)^{d_X/2} |\Sigma_X|^{1/2}} \exp \left\{ -\frac{(x - \mu_X) [\Sigma_X]^{-1} (x - \mu_X)^T}{2} \right\}.$$

## Joint multivariate normal vectors

$X$  and  $Y$  are *joint multivariate normal* if the combined vector

$$Z = \begin{pmatrix} X \\ Y \end{pmatrix} \quad (24)$$

is multivariate normal. In this case, we write

$$\mu_Z = \begin{pmatrix} \mu_X \\ \mu_Y \end{pmatrix}, \quad \Sigma_Z = \begin{pmatrix} \Sigma_X & \Sigma_{XY} \\ \Sigma_{YX} & \Sigma_Y \end{pmatrix}, \quad (25)$$

where

$$\Sigma_{XY} = \text{Cov}(X, Y) = \mathbb{E}[(X - \mu_X)(Y - \mu_Y)^T]. \quad (26)$$

- ▶ For joint multivariate normal random variables  $X$  and  $Y$ , we have the useful property that the conditional distribution of  $X$  given  $Y = y$  is multivariate normal, with conditional mean and variance

$$\begin{aligned} [\text{KF1}] \quad \mu_{X|Y}(y) &= \mu_X + \Sigma_{XY} \Sigma_Y^{-1} (y - \mu_Y), \\ [\text{KF2}] \quad \Sigma_{X|Y} &= \Sigma_X - \Sigma_{XY} \Sigma_Y^{-1} \Sigma_{YX}. \end{aligned}$$

- ▶ We write this as

$$X | Y = y \sim N[\mu_{X|Y}(y), \Sigma_{X|Y}]. \quad (27)$$

- ▶ The joint multivariate normal has a special property that the conditional variance of  $X$  given  $Y = y$  does not depend on the value of  $y$ . In non-Gaussian situations, it will usually depend on  $y$ .
- ▶ If  $\Sigma_Y$  is not invertible, we can interpret  $\Sigma_Y^{-1}$  as a generalized inverse.

## Notation for the Kalman filter recursions

We define the conditional means and variances for the filtering, prediction and smoothing distributions:

$$[\text{KF3}] \quad X_n | Y_{1:n-1} = y_{1:n-1} \sim N[\mu_n^P(y_{1:n-1}), \Sigma_n^P],$$

$$[\text{KF4}] \quad X_n | Y_{1:n} = y_{1:n} \sim N[\mu_n^F(y_{1:n}), \Sigma_n^F],$$

$$[\text{KF5}] \quad X_n | Y_{1:N} = y_{1:N} \sim N[\mu_n^S(y_{1:N}), \Sigma_n^S].$$

- ▶ For data  $y_{1:N}$ , we call  $\mu_n^P = \mu_n^P(y_{1:n-1}) = \mathbb{E}[X_n | Y_{1:n-1} = y_{1:n-1}]$  the *prediction mean*, and  $\Sigma_n^P$  the *prediction variance*.
- ▶  $\mu_n^F = \mu_n^F(y_{1:n}) = \mathbb{E}[X_n | Y_{1:n} = y_{1:n}]$  is the *filter mean* and  $\Sigma_n^F$  the *filter variance*.
- ▶  $\mu_n^S = \mu_n^S(y_{1:N}) = \mathbb{E}[X_n | Y_{1:N} = y_{1:N}]$  is the *smoothed mean* and  $\Sigma_n^S$  the *smoothed variance*.

## The Kalman matrix recursions

- ▶ Applying the properties of linear combinations of Normal random variables, we get the Kalman filter and prediction recursions:

$$[\text{KF6}] \quad \mu_{n+1}^P(y_{1:n}) = \mathbf{A}_{n+1}\mu_n^F(y_{1:n})$$

$$[\text{KF7}] \quad \Sigma_{n+1}^P = \mathbf{A}_{n+1}\Sigma_n^F\mathbf{A}_{n+1}^T + \Sigma_{X,n+1},$$

$$[\text{KF8}] \quad \Sigma_n^F = ([\Sigma_n^P]^{-1} + \mathbf{B}_n^T\Sigma_{Y,n}^{-1}\mathbf{B}_n)^{-1},$$

$$[\text{KF9}] \quad \mu_n^F(y_{1:n}) = \mu_n^P(y_{1:n-1}) + \Sigma_n^F\mathbf{B}_n^T\Sigma_{Y,n}^{-1}\{y_n - \mathbf{B}_n\mu_n^P(y_{1:n-1})\}$$

## Outline of a derivation of the Kalman matrix recursions

- ▶ The prediction recursions [KF6] and [KF7] follow from the property that if  $X$  is a  $d$ -dimensional multivariate normal,  $X \sim N(\mu, \Sigma)$ , then  $\mathbf{A}X + b \sim N(\mathbf{A}\mu + b, \mathbf{A}\Sigma\mathbf{A}^T)$ .
- ▶ Note that the multivariate normal identities [KF1,KF2] also hold when all variables are conditioned on some additional joint Gaussian variable, in this case  $Y_{1:n-1}$ .
- ▶ [KF8] and [KF9] can be deduced by writing out the joint density,

$$f_{X_n Y_n | Y_{1:n-1}}(x_n, y_n | y_{1:n-1}) \quad (28)$$

and completing the square in the exponent. The conditional density of  $X_n$  given  $Y_{1:n}$  is proportional to this joint density, with proportionality constant allowing integration to one.

**Exercise.** The derivation of the Kalman filter is not central to this course. However, working through the algebra to your own satisfaction is a good exercise.

- ▶ The Kalman filter matrix equations are easy to code, and quick to compute unless the dimension of the latent space is very large.
- ▶ In numerical weather forecasting, with careful programming, they are solved with latent variables having dimension  $d_X \approx 10^7$ .
- ▶ A similar computation gives backward Kalman recursions. Putting the forward and backward Kalman recursions together, as in [P9], is called *Kalman smoothing*.

## Further reading

- ▶ The approach in this chapter is aligned with King, Nguyen, and Ionides [1].
- ▶ Chapter 6 of Shumway and Stoffer [2] gives an approach emphasizing linear Gaussian state space models.

## Acknowledgements

- ▶ Compiled on February 25, 2026 using Python.
- ▶ Licensed under the Creative Commons Attribution-NonCommercial license. Please share and remix non-commercially, mentioning its origin.
- ▶ We acknowledge previous versions of this course.

## References I

- [1] Aaron A. King, Dao Nguyen, and Edward L. Ionides.  
“Statistical Inference for Partially Observed Markov Processes via the R Package *pomp*”. In: *Journal of Statistical Software* 69.12 (Mar. 2016), pp. 1–43. DOI: 10.18637/jss.v069.i12.
- [2] Robert H Shumway and David S Stoffer. *Time Series Analysis and its Applications: With R Examples*. 4th. Springer, 2017.